

# SVHN Classification with CNN — Detailed Project Report

## A. Introduction

The objective of this project is to build a robust digit recognition model using the **Street View House Numbers (SVHN)** dataset, a real-world dataset composed of digit images from Google Street View. Unlike MNIST, SVHN images are in color and contain complex backgrounds, making this a more challenging and realistic classification problem. Each image is a cropped digit (0–9) from house numbers in street scenes, presented as a 32×32 RGB image.

Applications include:

- - Address and street sign recognition
- - Automated postal sorting
- - OCR in navigation and mapping apps
- - Smart cameras and intelligent transportation systems

## B. Dataset Overview

- - **Source:** SVHN (Cropped Digits) from `.mat` files
- - **Train Samples:** ~73,000
- - **Test Samples:** ~26,000
- - **Classes:** 10 digits (0–9); label '10' is mapped to '0'
- - **Format:** 32×32 RGB images; cropped digits

Key characteristics:

- - **Complex backgrounds:** Unlike MNIST, SVHN includes background clutter
- - **Color variance:** Input is 3-channel (RGB), not grayscale
- - **Real-world noise:** Some images include neighboring digits or lighting issues

## C. Methodology (Detailed)

### ### 1. Data Loading

Data is loaded from MATLAB `.mat` files using `scipy.io.loadmat`. Each file contains:

- - ``X``: 4D image tensor with shape (32, 32, 3, N)
- - ``y``: Label vector with shape (N,)

To prepare this data:

- - Transpose ``X`` to shape (N, 32, 32, 3)
- - Normalize pixel values to [0, 1] by dividing by 255
- - Replace label '10' with '0'

### ### 2. Preprocessing

After loading, preprocessing includes:

**\*\*a. One-hot encoding:\*\***

- - Converts numeric labels (0–9) to categorical vectors for softmax output layer compatibility.

**\*\*b. Train-validation split:\*\***

- - To prevent overfitting, 20% of training data is used for validation using ``train_test_split``.

**c. Data augmentation:** **(optional but beneficial)**

Using ``ImageDataGenerator``, the dataset is artificially expanded with real-time augmentation:

- - Rotation ( $\pm 10^\circ$ )
- - Zoom ( $\pm 10\%$ )
- - Horizontal and vertical shifts ( $\pm 10\%$ )

This helps the model generalize better on unseen data by exposing it to slight distortions.

### ### 3. CNN Model Architecture

A custom Convolutional Neural Network (CNN) is built using ``Sequential`` API. It includes:

- - **Convolutional layers (Conv2D):**
  - - Extract spatial features from the image
  - - Filters of sizes 32, 64, and 128 progressively learn low to high-level features
- - **BatchNormalization:**
  - - Standardizes outputs of layers to stabilize and speed up training
- - **MaxPooling2D:**

- - Downsamples the feature maps, reducing spatial dimensions
- - **Dropout:**
- - Randomly disables neurons to prevent overfitting
- - **Dense (Fully Connected) layers:**
- - Translates feature maps into class probabilities
- - **Softmax output:**
- - Produces a probability distribution over 10 classes

Optimizer: **Adam** (adaptive learning rate)

Loss Function: **Categorical Crossentropy**

### ### 4. Training Strategy

- - **Batch size:** 128
- - **Epochs:** 50 (with early stopping)
- - **Callbacks:**
- - `EarlyStopping`: Stops training if validation loss doesn't improve
- - `ReduceLROnPlateau`: Reduces learning rate on performance plateau

Training is monitored for both loss and accuracy on training and validation sets. This allows tracking of underfitting or overfitting behavior.

## D. Model Evaluation

### ### 1. Accuracy

Model achieved ~96% test accuracy — a strong result considering the complexity of SVHN.

### ### 2. Confusion Matrix

A 10x10 matrix showing true labels vs. predicted labels, used to identify:

- - Which digits are misclassified most
- - Confusion trends (e.g. 3 vs. 5, 8 vs. 0)

### ### 3. Classification Report

Includes:

- - **Precision**: Correct positive predictions / total predicted positives
- - **Recall**: Correct positive predictions / total actual positives
- - **F1-score**: Harmonic mean of precision and recall

### ### 4. Misclassified Examples

10 incorrectly predicted images are visualized. Most errors are due to:

- - Blurred or partially occluded digits
- - Side digits not removed completely during cropping
- - Background color blending with digit

## E. Results Summary

Digit-wise Precision, Recall, F1-Score:

0: 96%, 97%, 96%

1: 98%, 99%, 99%

2: 95%, 94%, 95%

3: 94%, 92%, 93%

4: 96%, 95%, 95%

5: 93%, 92%, 93%

6: 95%, 96%, 96%

7: 96%, 95%, 96%

8: 94%, 95%, 95%

9: 94%, 93%, 94%

**Macro average:** Precision = 0.95, Recall = 0.95, F1 = 0.95

## F. Challenges and Limitations

- - **Cluttered backgrounds**: Some cropped images still contain side digits
- - **Digit similarity**: 3 vs 5, 8 vs 0 misclassifications
- - **Small variations**: Rotations and occlusions affect predictions

## G. Future Improvements

1. **Use the Extra Set**: SVHN provides ~500k additional labeled samples.

2. **\*\*Switch to Transfer Learning\*\***: Use EfficientNetB0, MobileNet, or ResNet with pretrained weights.
3. **\*\*CTC-based digit string recognition\*\***: Instead of cropped digits, build models that detect and recognize multi-digit strings.
4. **\*\*Hyperparameter tuning\*\***: Explore optimizers, dropout rates, and learning rate schedules.
5. **\*\*Ensemble models\*\***: Combine predictions of multiple CNNs.

## H. Conclusion

This project demonstrates that a well-designed CNN can perform strongly on real-world digit recognition using the SVHN dataset. The model handles noise and complexity well, achieving ~96% accuracy. With more data and fine-tuning, it could be deployed in real applications like address recognition and smart OCR.